

CS211: Computer Architecture

Summer 2020

Instructors and TAs

- Instructor:
 - David Domingo
 - Email: david.domingo@rutgers.edu
 - Office Hours: Fridays 2pm-3pm
- Teaching Assistants:
 - Yujie Ren
 - Email: yujie.ren@rutgers.edu
 - Office Hours: Mondays 1pm-2pm
 - Shaleen Garg
 - Email: shaleen.garg@rutgers.edu
 - Office Hours: Thursdays 11am-12pm
 - Chengguizi Han
 - Email: chengguizi.han@rutgers.edu
 - Office Hours: Thursdays 2pm-3pm

Course Logistics

- Sakai will be the primary source of course information
 - <https://sakai.rutgers.edu>
 - Announcements, course material, assignments, etc.
- Lectures
 - Held synchronously online via Sakai Meetings
 - Slides will be uploaded after lecture
 - Lectures will be recorded and available*
- Office Hours
 - Held online via Sakai Meetings
- Piazza
 - Main forum for questions and discussion on lecture materials and assignments.

Textbooks

- Required
 - **Computer Systems:A Programmer's Perspective,** by R.E.Bryant and D.R.O'Hallaron
- Recommended:
 - **Computer Organization and Design:The Hardware/Software Interface** by D.A.Patterson and J.L.Hennessy
 - **The C Programming Language** by B.W.Kernighan and D.M.Ritchie
 - Any book about C similar to
(https://publications.gbdirect.co.uk/c_book/)

What You Should Know

- Prerequisite: CS 112 Data Structures
- Programming Languages
- Algorithms
- Data Structures
- Basics of how to write, run, and test programs

Goal

- Understand how programs run on hardware
- Write programs with good performance on modern architecture computers.

What you will learn

- How to program in C and Assembly
- The major hardware components in computer systems
- Trends in technology and computer architecture
- How hardware components are built from digital logic
- How programs written in a high-level language (e.g., C) is actually executed by the hardware
- How to understand and improve the performance of programs

Course Expectations

- Fun part:
 - 3+ Projects (programming assignments)
- Not So Fun Part:
 - 1 Midterm and 1 Final*
- What we expect from you
 - Attend Lectures
 - Ask Questions
 - Ask for help if you feel lost
 - Start programming assignments early
 - Do not copy or cheat

No Late Assignments

- We will not accept late assignments
- Emails with assignment attach will be discarded
- Assignments must be handed-in on Sakai
 - Deadline enforces by Sakai
 - Can submit unlimited number of times

Collaboration vs. Cheating

- Collaboration is encouraged!
 - Learn by discussing and helping each other
 - But, you must not cheat and copy
- Cheating will not be tolerated
 - We will look for cheating
 - Once, found everyone involved will be punished
 - <https://www.cs.rutgers.edu/academics/undergraduate/academic-integrity-policy>
- If you having any trouble with the course seek help
 - Email me or the TAs

List of topics

- Hardware trends
- C programming
- Data representation
- Assembly program
- Memory Hierarchy/Caching
- Digital Logic

Programming Assignments

- 3 Programming Assignments
- Program in C and/or Assembly
 - Start early, do not wait until the last minute
- Programming and Grading done using iLab Machines
 - <https://resources.cs.rutgers.edu/docs/instructional-lab/>
 - Sign up for account if you haven't already
<https://services.cs.rutgers.edu/accounts/>
 - Programming in Linux Environment

Grading

- **Grading:**
 - 45%: Programming Assignments
 - 20%: Midterm
 - 30%: Final Exam
 - 5%: Class Participation
 - Asking Questions in Class
 - Asking/answering questions in Piazza
 - All exams are cumulative
 - No make-up exams except for university sanctioned reasons
 - Must inform me before the exam

Tentative Schedule

- June 22nd – July 1st : C Programming
- July 6th : Data Representation
- July 8th – July 20th : Assembly Language
- July 22nd : Midterm
- July 29 : Memory Hierarchy/Cache
- July 31st – August 10 : Digital Logic
- August 12th : Final Exam

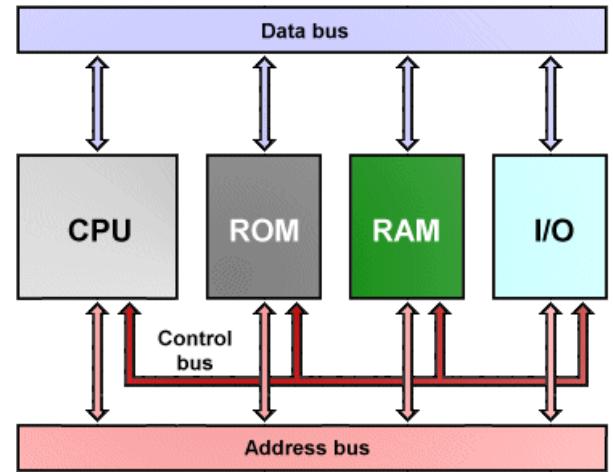
Computer Architecture: An Introduction

Computer Architecture



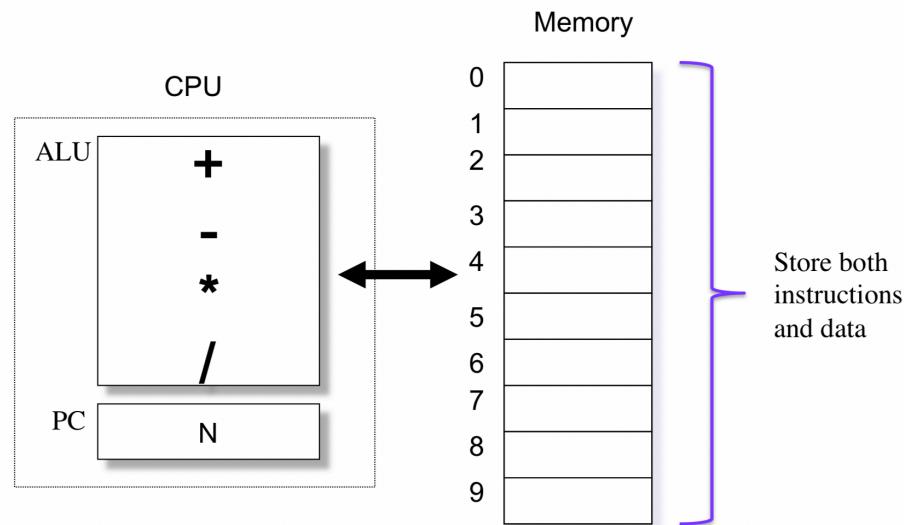
Main Components

- CPU
 - Central Processing Unit
 - Performs computation
- Memory
 - Stores data
- BUS
 - Used for carrying information
 - Communication and data transfer
- I/O Devices
 - Human Interface (mouse, keyboard, screen)
 - Networking
 - Graphics
 - Storage

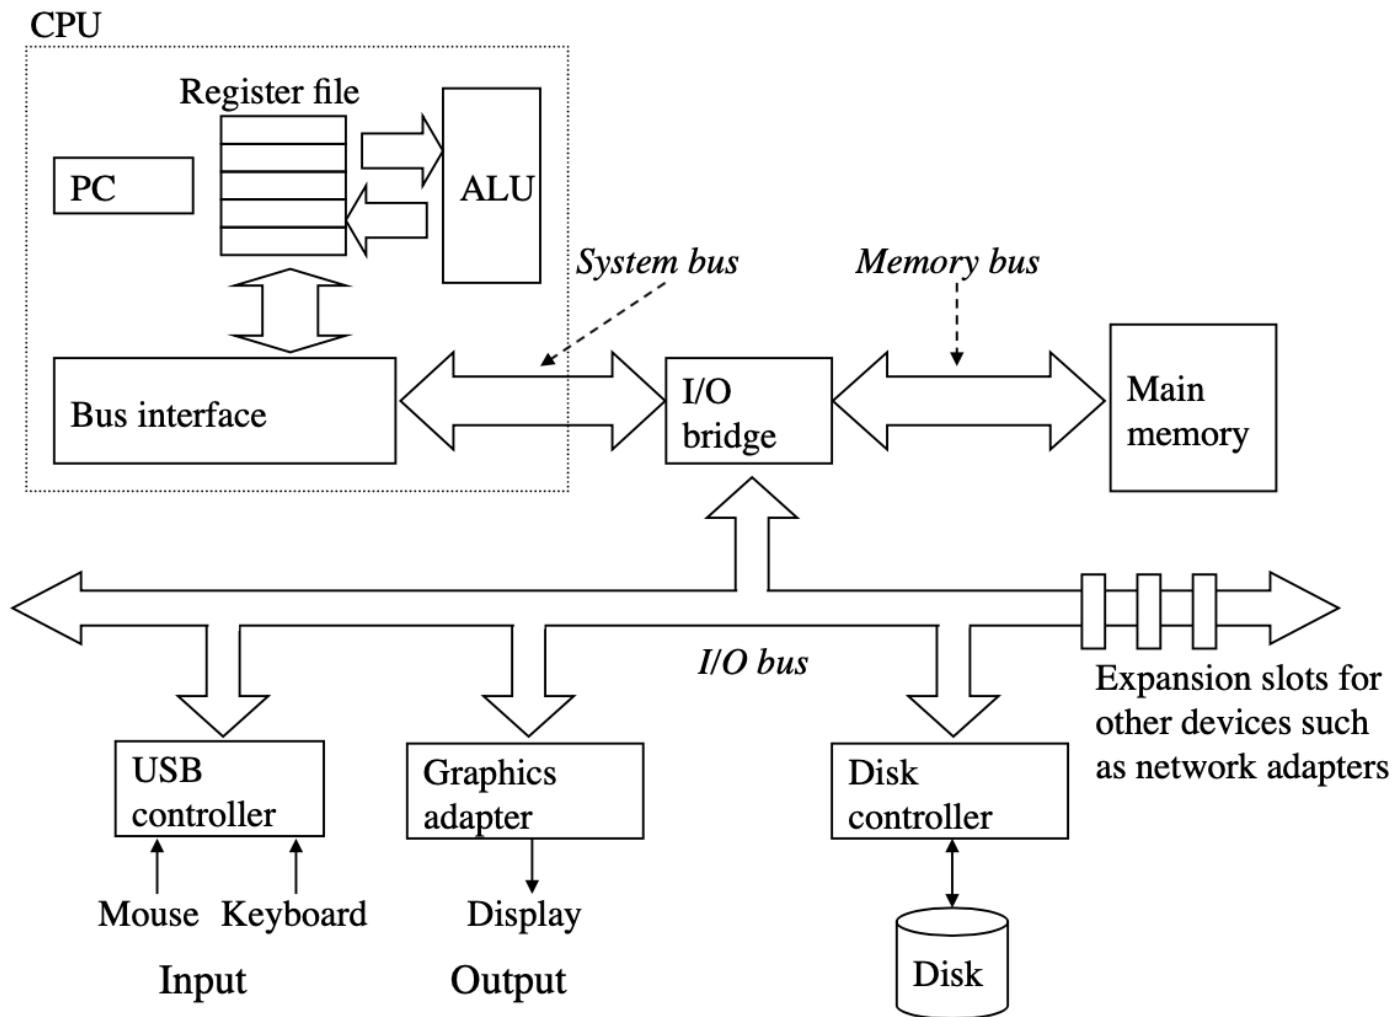


Von Neuman Model

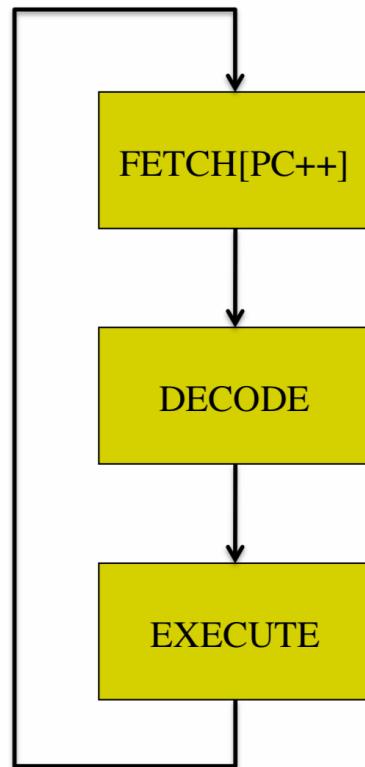
- Computer Architecture Model
- Central processing unit
 - Arithmetic Processing Unit (ALU)
 - Control Unit (instruction register and program Counter)
- Memory Component that stores both data and instructions



Von Neumann Model: In Practice



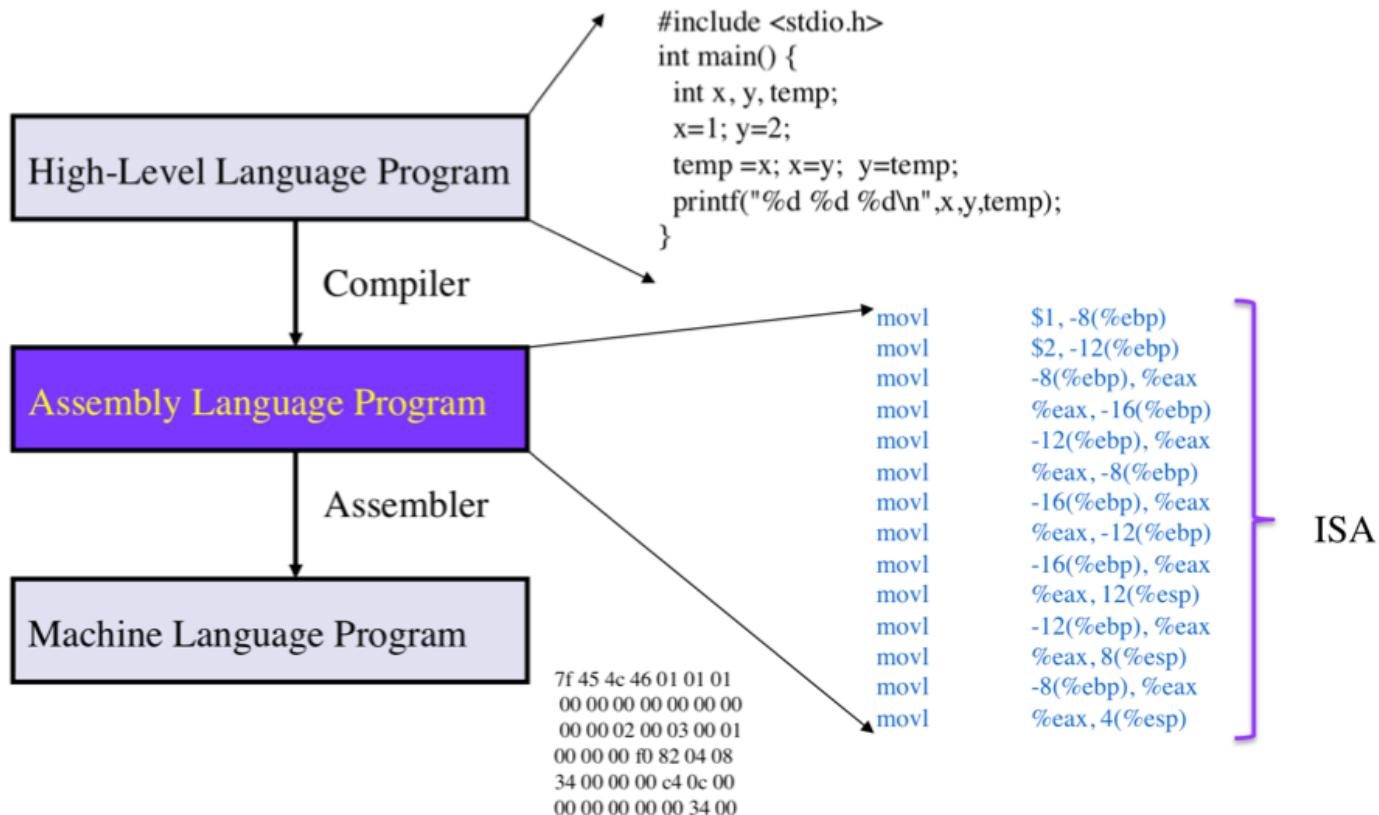
Basic CPU Function



Arithmetic: +, -, *, /
Logic: bre, jmp

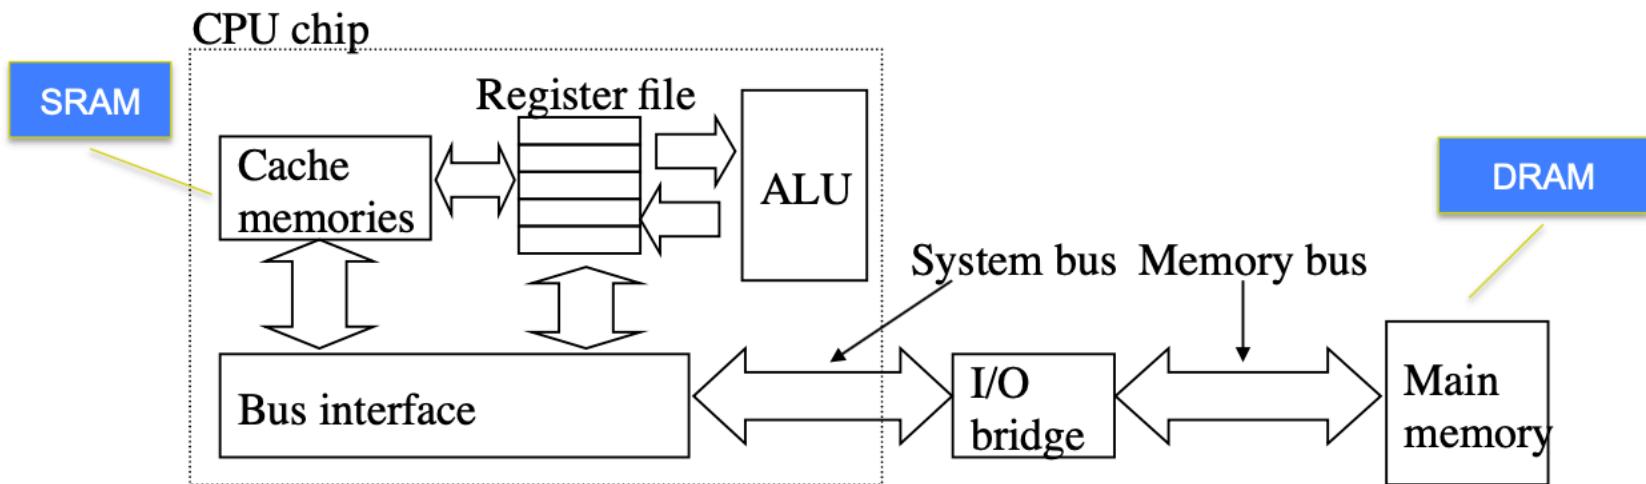
Program to Hardware

- How logic turns into something machines can understand



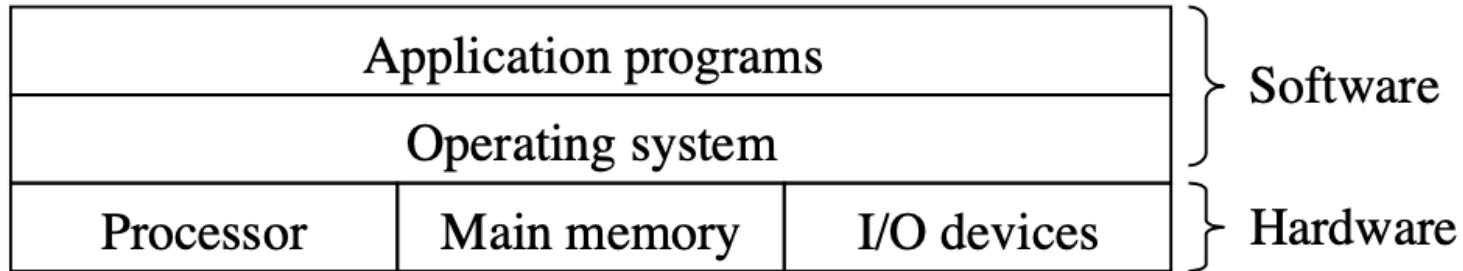
Memory Hierarchy

- A program that resides on disk must be loaded into memory for execution
- Why do we have levels of memory?
 - Larger devices are slower
 - Faster devices are more expensive
 - Typically less storage per dollar

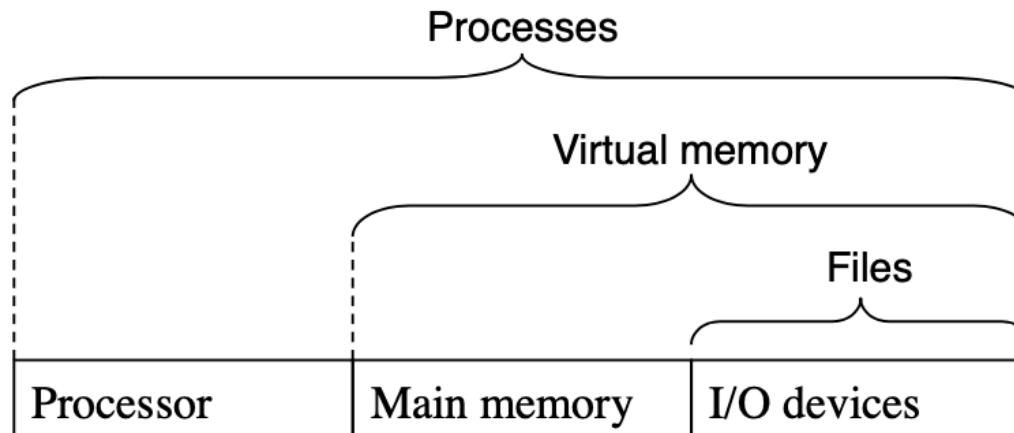


The Operating System

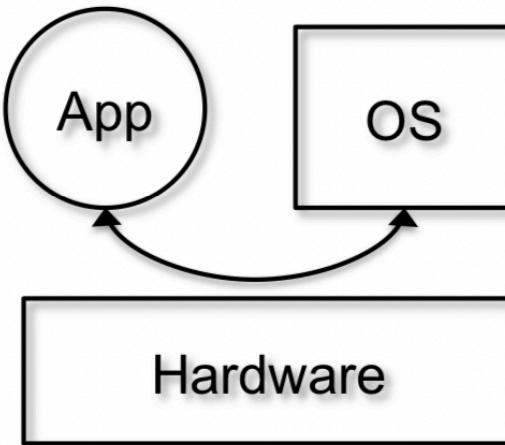
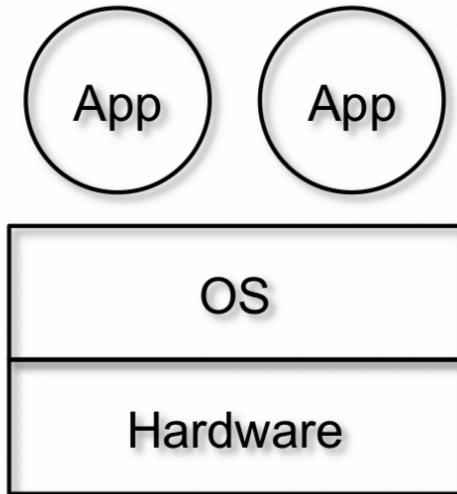
- An operating system manages the hardware



- Manages hardware through abstractions



Operating System Reality



- Main Idea: OS sits on top of hardware facilitate applications applications
- Reality: Apps typically run directly on hardware. Will switch over when OS when needed

Computer Architecture

- How to design and build components
- In this class
 - Understand how programs run on current system
 - Understand how current architecture affect programming languages
- Additionally you'll be able to think about
 - How can I make programs faster
 - How can I make programs more power efficient

Architecture Trends

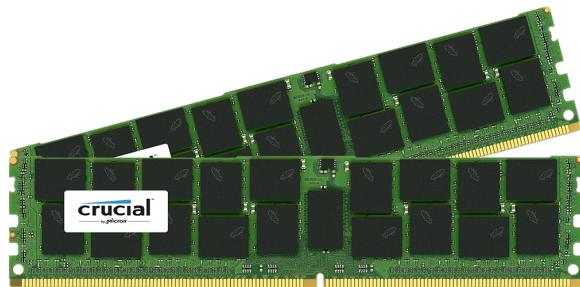
- How have Computer Architecture changed over time?



CPU



Storage



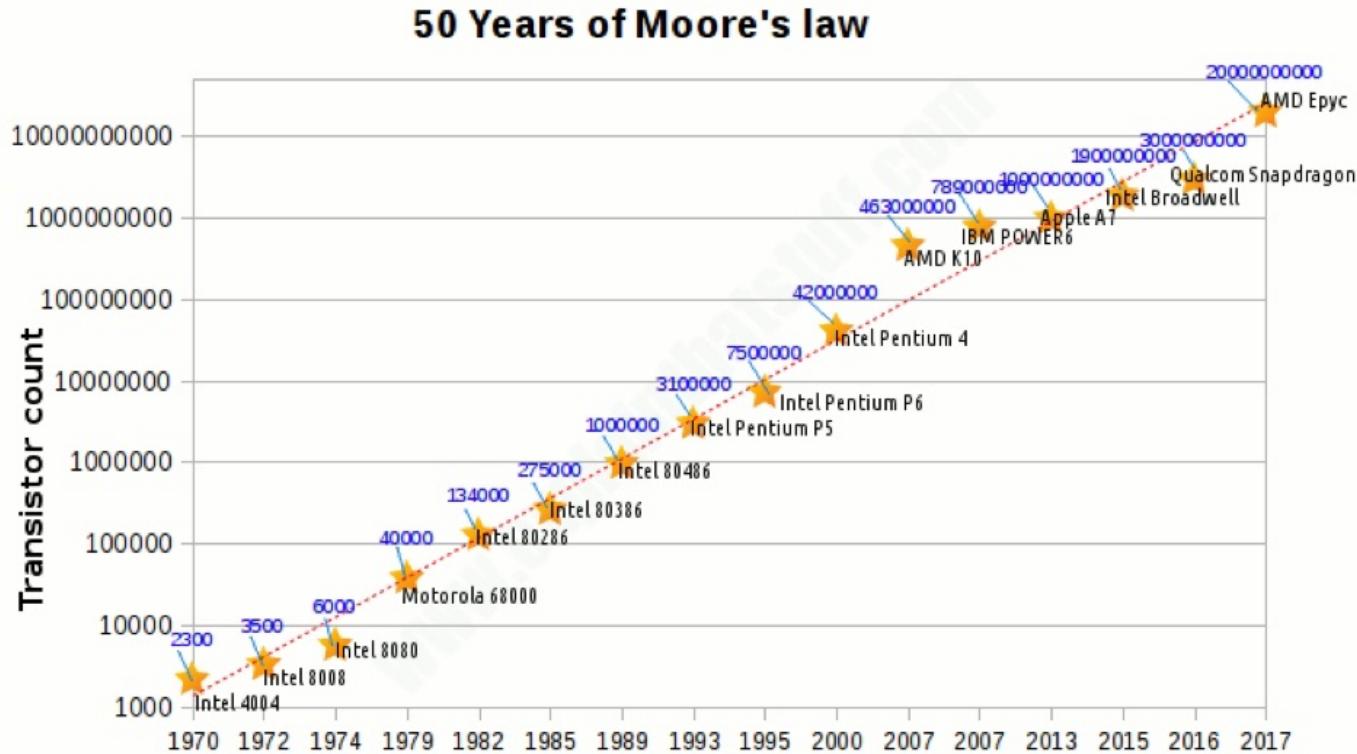
Memory

Moore's Law

- Gordon Moore, an Intel Engineer
- Observed architectural improvement trend
 - # of transistors on a chip doubles every 18 months
- General exponential growth in hardware:
 - Process speed x2 every 18 months
 - Memory capacity x2 every 2 years
 - Disk capacity x2 every year

Transistor Count Trend

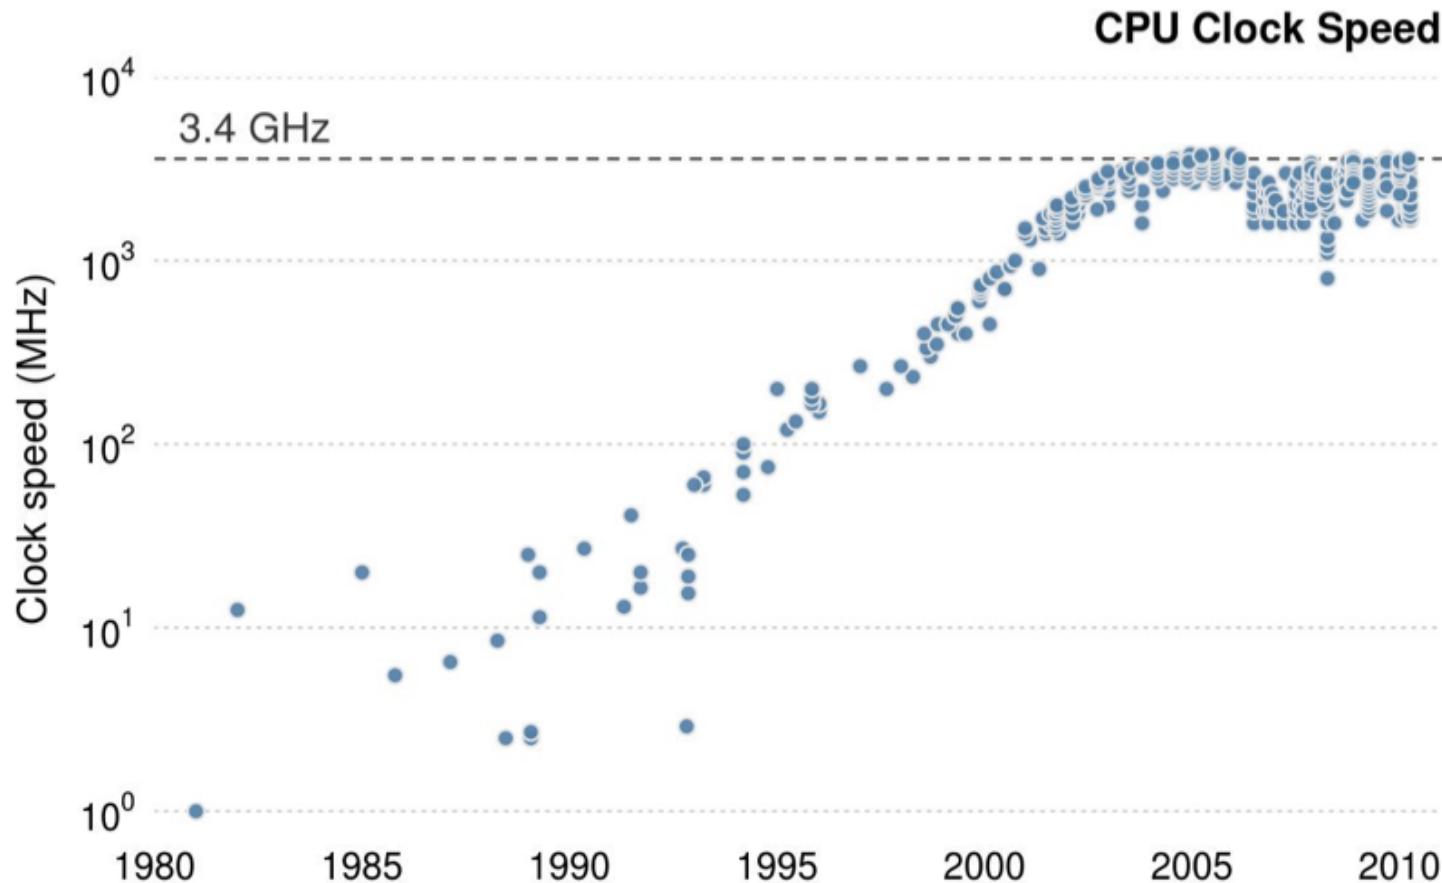
- Transistors are needed to implement logical gates within a CPU



- 32-core AMD Epyc has 19,200,000,000 transistors

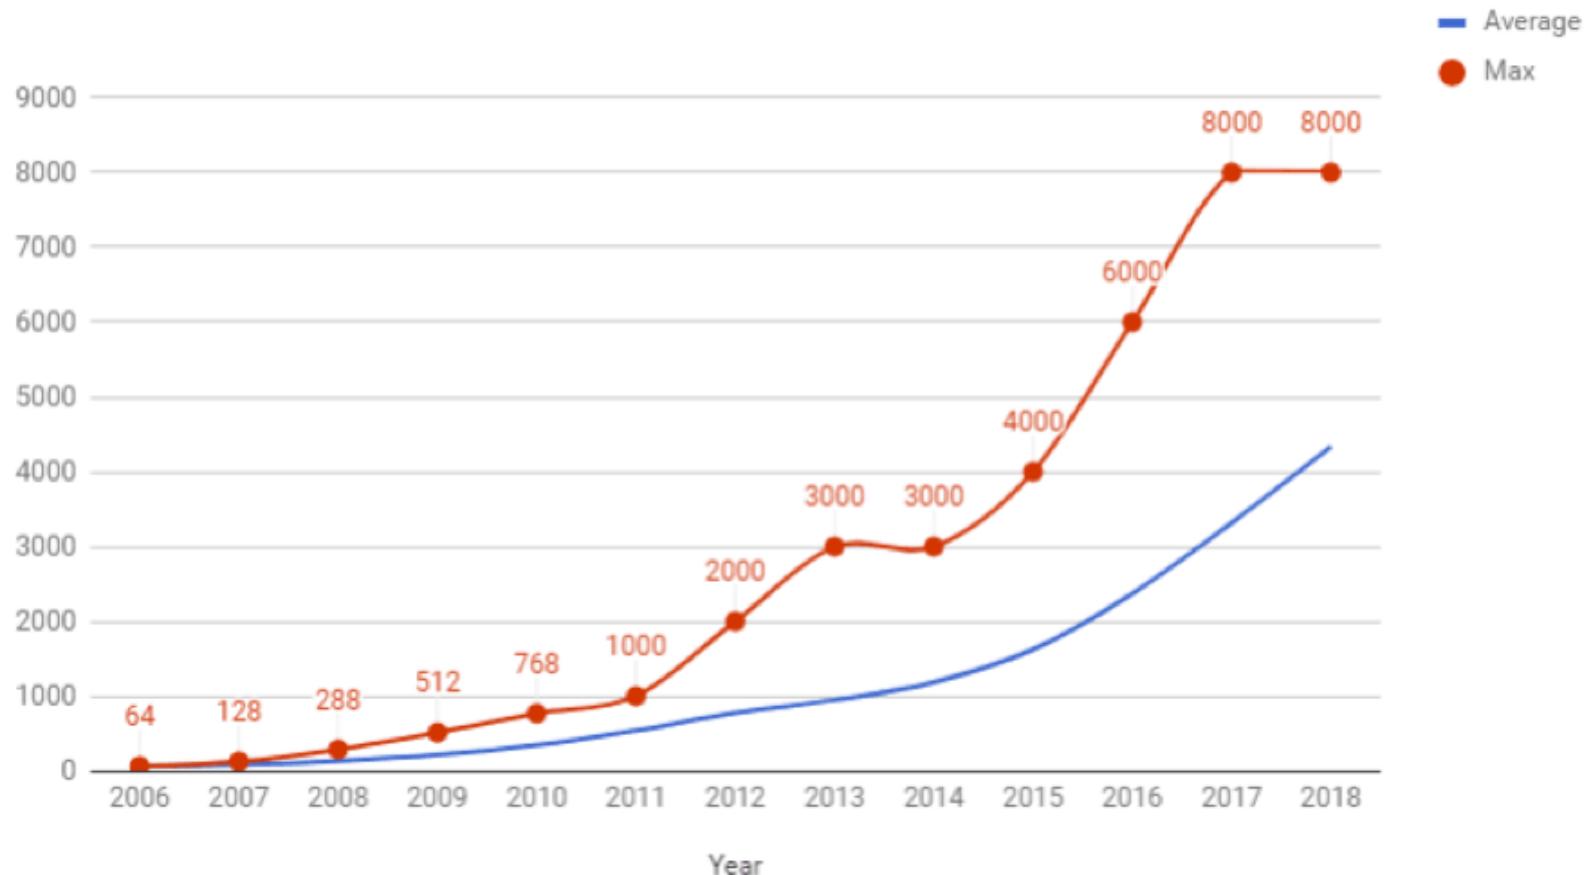
Trend: Clock Speed

- Clock speed is the rate at which a processor completes a processing cycle



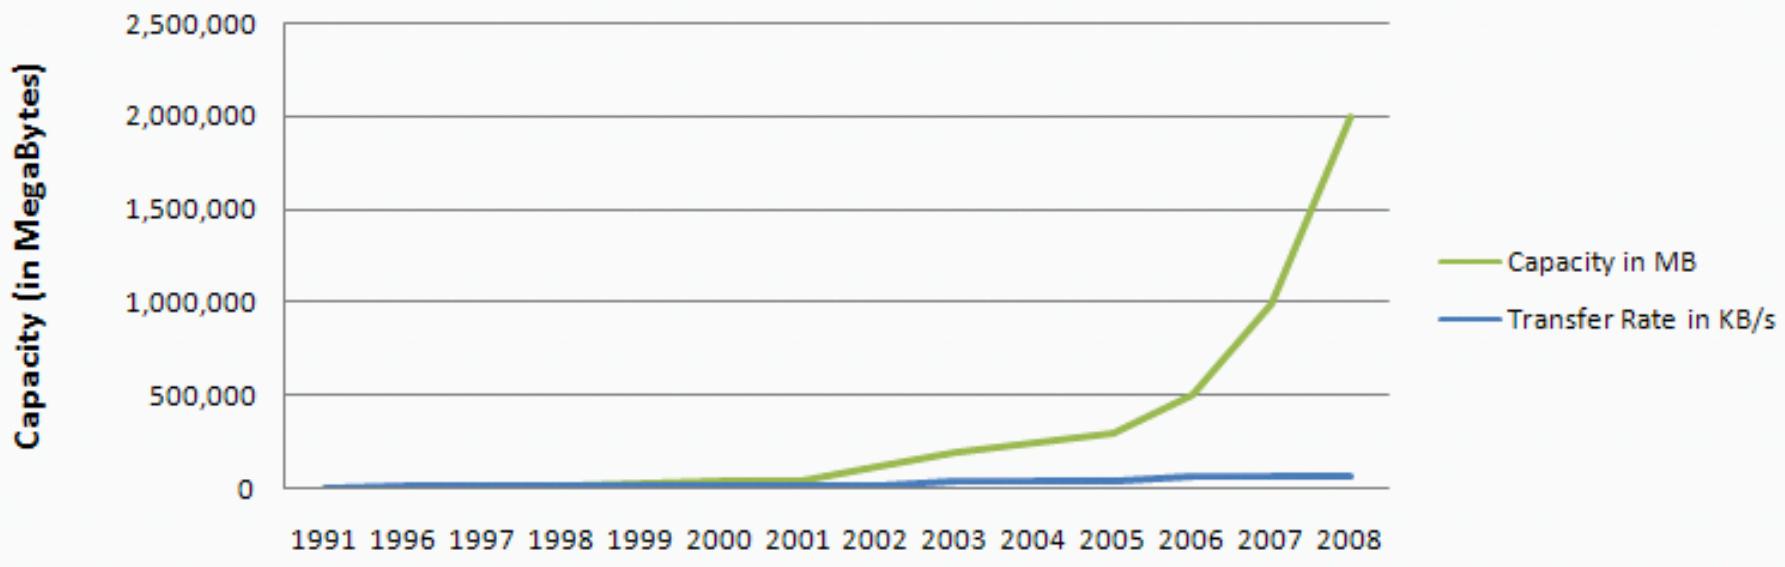
Trend: Memory Capacity

RAM capacity through the years (in MB)



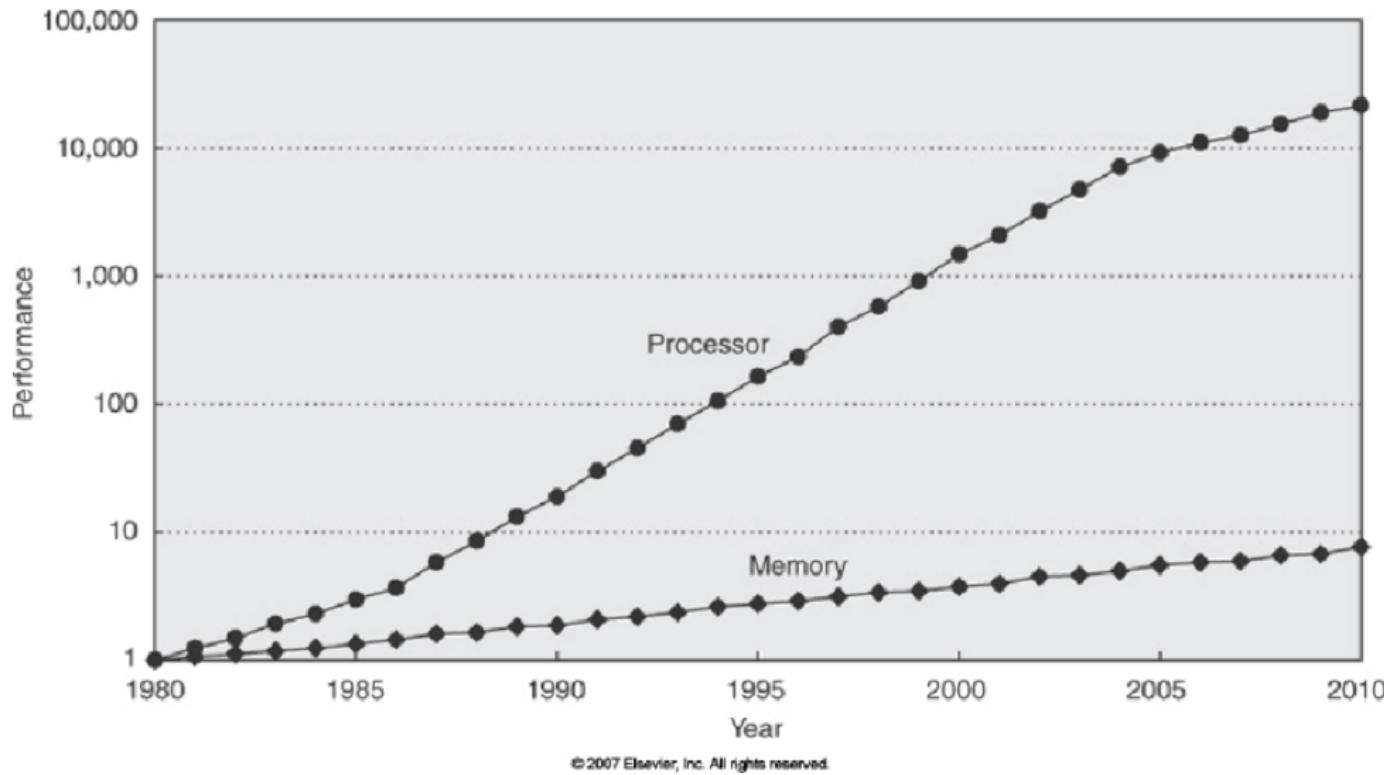
Trend: Disk Capacity

Relative Improvement
Hard Disk Capacity v.s. Disk Transfer Performance



Why are these trends important when it comes to designing system components?

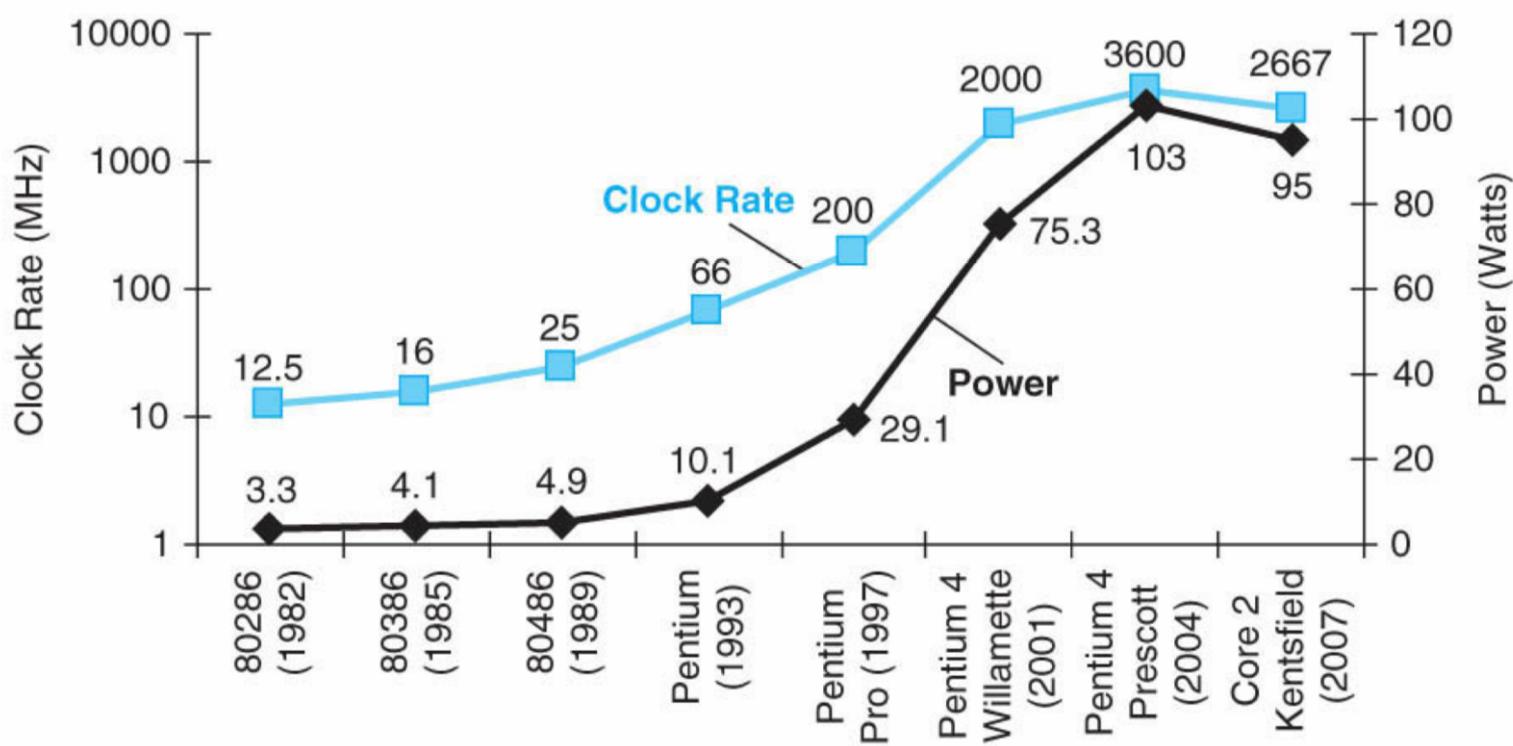
CPU vs Memory Performance Gap



- CPU performance grows faster than Memory
- Implication:
 - Need to make efficient use of the memory hierarchy
 - Registers/Caches

Power Wall

- More transistors -> higher power densities -> higher temperature
- Scaling performance limited by fundamental constraints
 - Power delivery
 - Cooling
- Goal is to maximize performance per watt



In Summary

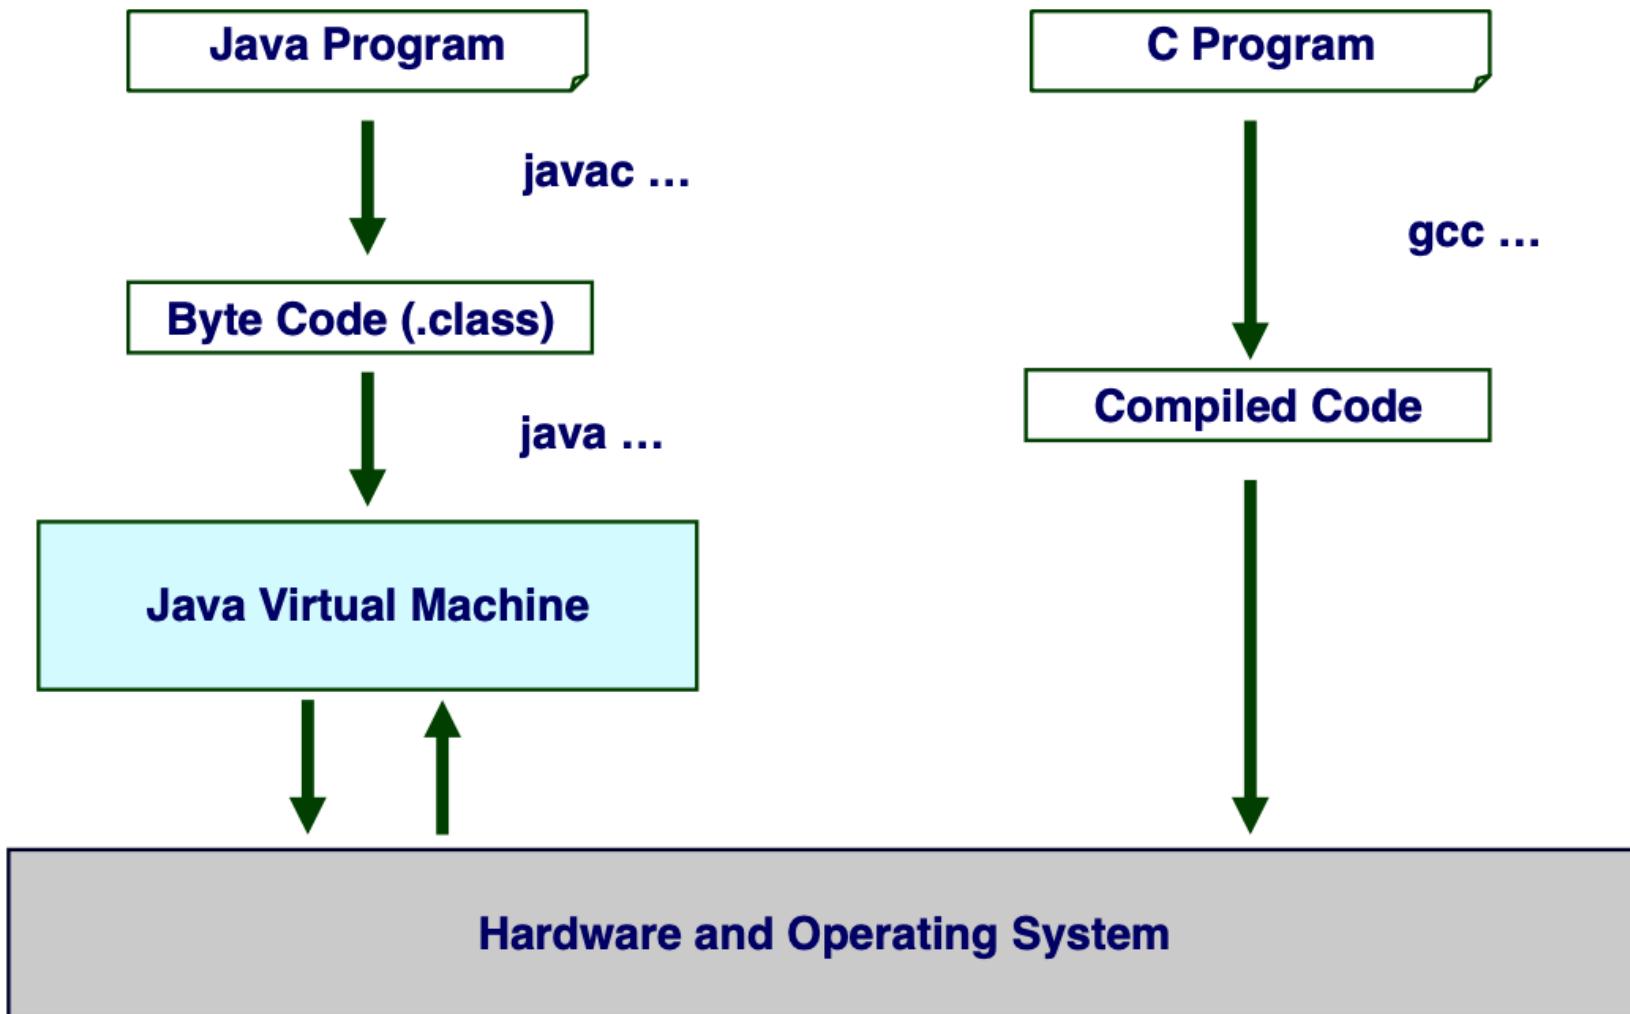
- Todays systems are complex
- Faster processor and systems enable new domains of applications
- Understanding the system is crucial to
 - Develop efficient applications
 - Ensure good tradeoffs between performance/power etc.
 - Address various constraints (memory, power, etc.)

Introduction to C Programming

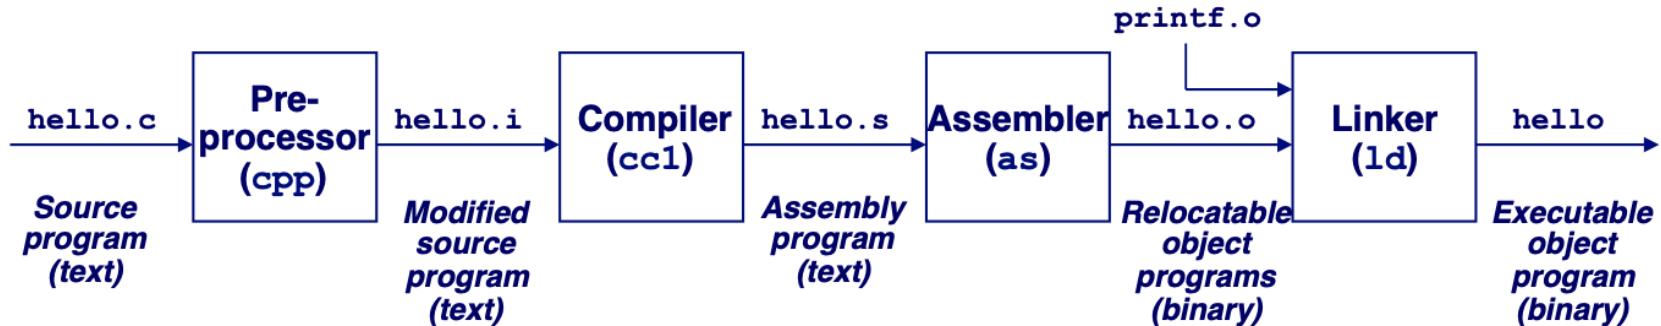
Why Learn C?

- You are learning to be a computer scientist
 - Languages are just tools
 - Choose the tools appropriate for the task
- Current task: Learn how programs written in high-level languages run on computers
- C is closer to machine compared to higher level languages such as Java

Comparison with Java



C Program to Machine Execution



- **Compilation System Phases:**
 1. Preprocessing: process directives
 2. Compilation: translated into assembly code
 3. Assembly: translates into machine language
 4. Linking: merges libraries

C vs Java Comparison

C

- Procedural
- Compiled
- Low-Level
- User manages memory

Java

- Object Oriented
- Interpreted
- High-Level
- Garbage collector

Anatomy of a C Program

```
#include <stdio.h>
#include <stdlib.h>
```

include files

```
char cMessage[] = "Hello\n";
```

declaration of global variables

```
/* Execution will start here */
int main (int argc, char **argv)
{
```

comment

```
    int i, count;
    count = atoi(argv[1]);
    for (i = 0; i < count; i++) {
        printf("Hello %d\n", i);
    }
}
```

one or more function;
each program starts execution at “main”

declaration of local variables

code implementing function

Commenting

- Two ways to comment
- Start comment with //
 - comments until end of line
 - Example:

```
// This is a comment
// This is a second comment
```
- Start comment with /* and end with */
 - comments can span multiple lines
 - Example:

```
/* This is a comment
that spans multiple lines */
```
- Comments are critical for good development

Variable Declarations

- Each variable has a type, which tells how the compiler how the data is to be interpreted (and how much space it needs, etc)
 - Examples:
int counter;
char letter;
- Variables can be global or local
 - Global : declare outside scope of any function and can be accessed anywhere
 - Local: declare inside the scope of a function, only accessible from inside of the function

Basic Data Types

Keyword	Data Type	Examples
char	individual character	'a', 'b', '\t', '\n'
int	integers	-10, 15, 0, 1324
float	real numbers	-24.4, 0, 4.23
double	real numbers with double precision	-24.4, 0, 4.23

- **Modifiers**

- **short, long** : control size/range of numbers
- **signed, unsigned**: include negative numbers or not

Arithmetic Operators

Symbol	Operation	Usage
+	addition	X + Y
-	subtraction	X - Y
*	multiplication	X * Y
/	division	X / Y
%	modulus	X % Y

Special Operators

- Changes value of variable before (or after) its value is used in an expression
- `++` increments a variable
- `--` decrements a variable
- Can be used before or after a value

Use	Operation
<code>X++</code>	post increment
<code>++X</code>	pre increment
<code>X--</code>	post decrement
<code>--X</code>	pre decrement

- pre: increment/decrement before using its value
- post: increment/decrement after using its value

Relational Operators

Symbol	Operation	Usage
>	greater than	$X > Y$
\geq	greater or equal to	$X \geq Y$
<	less than	$X < Y$
\leq	less than or equal to	$X \leq Y$
\equiv	equal	$X \equiv Y$
\neq	not equal	$X \neq Y$

- Result is 1 (true) or 0 (false)
- Remember not to confuse equality (\equiv) with assignment ($=$)

Logical Operators

Symbol	Operation	Usage
!	logical NOT	$\text{!}X$
$\&\&$	logical AND	$X \&\& Y$
$\ $	logical OR	$X \ Y$

- Treats entire variable (or value) as True or False
- Result is either 0 (false) or non-zero (true)

Bit-wise Operators

Symbol	Operation	Usage
\sim	complement	$\sim X$
$\&$	bit-wise AND	$X \& Y$
$ $	bit-wise OR	$X Y$
$>>$	bit-wise shift right	$X >> 2$
$<<$	bit-wise shift left	$X << 2$

- Operate on bits of variables
- For example:
 - $\sim 0101 = 1010$
 - $0101 \& 1010 = 0000$
 - $0101 | 1010 = 1111$
 - $0101 >> 2 = 0001$
 - $0101 << 3 = 1000$

Control Statements

- Conditional
 - if else
 - Switch
- Iteration (loops)
 - for
 - while
 - do while
- Specialized “go-to”
 - break
 - continue

If Else Statements

```
if (expression){  
    statement  
}
```

```
if (expression A){  
    statement 1;  
} else {  
    statement 2;  
}
```

```
if (expression A){  
    statement 1;  
} else if (expression B){  
    statement 2;  
} else {  
    statement 3;  
}
```

- If an expression is true, then run the statement
- If an expression is true, then run the statement 1, else run another statement 2
- Evaluates all expressions until finds one with non-zero result
 - Executes corresponding statements
 - If all expressions evaluate to zero, executes statements for “else” branch

The Switch Statement

```
switch (expression) {  
    case const-1:  
        statement 1;  
    case const-2:  
        statement 2;  
    case const-3:  
        statement 3;  
    default:  
        statement-n;  
}
```

- Evaluates expression and skips to case corresponding to the result.
- Result must be an integer
- Executes statements corresponding to the result and continues executing until encountering a break or end of switch statements
- default always matches

Switch Statement (example)

```
int fork;  
...  
switch (fork) {  
    case 1:  
        printf("take left");  
    case 2:  
        printf("take right");  
        break;  
    case 3:  
        printf("make U turn");  
        break;  
    default:  
        printf("go straight");  
}
```

Iterations (Loops)

Format	Description
<pre>for (start-expression; condition; update-expression){ ... }</pre>	<ul style="list-style-type: none">- runs zero or more times- keeps iterating while cond-expression != 0- compute start-expr before first iteration- compute update-expr after each iteration
<pre>while (expression) { ... }</pre>	<ul style="list-style-type: none">- runs zero or more times- keeps iterating while expression != 0- compute expression before each iteration
<pre>do { ... } while (expression)</pre>	<ul style="list-style-type: none">- runs one or more times- keeps iterating while expression != 0- computes expression after each iteration

Specialized Go-to's

- `break;`
 - force immediate exit from switch or loop
 - goes to statement immediately following switch/loop
- `continue;`
 - skip the rest of the computation in the current iteration of the loop
 - goes to evaluation of conditionals expression for execution of the next iteration

Specialized Go-to's (example)

```
int index = 0;
```

```
int sum = 0;
```

```
while ( (index >= 0) && (index <= 20) ){
    index = index + 1;
    if (index == 11){
        break;
    }
    if ( (index % 2 ) == 1){
        continue;
    }
    sum = sum + index;
}
```

Functions

- Similar to Java Method

- Components

- Return Type
 - void if no return value
- Function name
- Parameters
- Body
 - statements to be executed
 - return forces exit from function, resuming at statement immediately after function call

```
int Factorial(int n) {  
    int i;  
    int result = 1;  
    for (i = 1; i <= n; i++) {  
        result = result * i;  
    }  
    return result;  
}
```

Function Calls

- Function call as part of an expression
 - Example: $z = x + \text{Factorial}(y);$
 - Arguments evaluated before function call
 - Return value is used to compute expression
 - Cannot have a void return type
- Function Call as a statement
 - Example: $\text{Factorial}(y);$
 - Can have a void return type
 - Returned value is discarded (if there is one)

Basic Input and Output

- Variety of I/O functions in C standard library
- Input:
 - `int printf(const char *format,);`
 - Writes to standard output based on the format provided
 - Examples:
 - `printf("Hello world\n");`
 - `printf("My number grade is %d\n", grade);`
- Input:
 - `int scanf(const char *format, ...);`
 - Reads input from standard input and scans it based on formatting provided
 - Examples:
 - `scanf("%d", &grade);`
 - `scanf("%d %c", &grade, &lettergrade)`
 - Note: Don't worry about the "&" yet

Basic Input and Output

Common format specifiers

Specifier	Type
%c	character
%d	integer
%u	unsigned int
%f	float
%lf	double
%s	string
%p	pointer

Common escape sequences

Escape Sequence	Character
\n	newline
\t	tab
\\\	backslash
\'	single quotation
\\"	double quotation

Basic C Program (HelloWorld.c)

```
#include <stdlib.h>
#include <stdio.h>

int main(int argc, char** argv){
    printf("Hello World!\n");
    return 0;
}
```

Compiling and Running a C Program

- Use gcc to compile your programs
 - Example: `gcc HelloWorld.c`
 - Will compile into executable named “`a.out`”
 - Run the program by running `./a.out`
- Use the “`-o`” flag to specify the output name
 - Example: `gcc HelloWorld.c -o HelloWorld`
 - Will compile program into executable named “`HelloWorld`”
 - Run the program by running `./HelloWorld`

For Next Lecture

- Start practicing C Programming
 - Connect to iLab Machines
 - Start coding, compiling, and running short C programs
 - Get used to the iLab Machine
- For those who want to get ahead
 - Arrays
 - Pointers
 - Structs

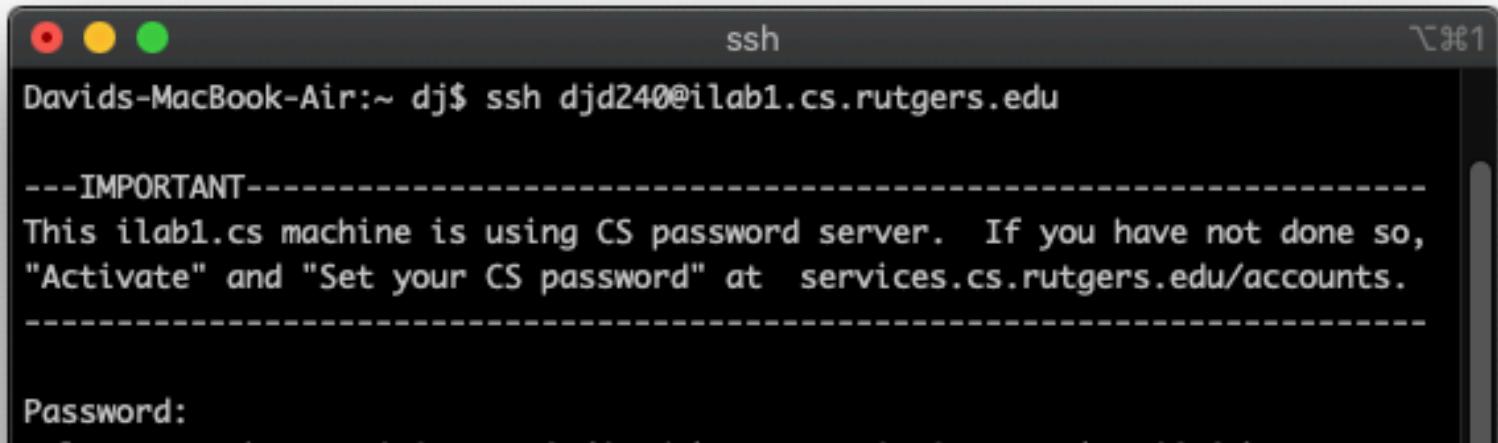
iLab Machines and Linux

iLab Machines

- Sign up for a iLab account if you haven't already
 - <https://services.cs.rutgers.edu/accounts/>
- Two Ways to use it
 - Via Secure Shell (SSH)
 - terminal (Mac/Linux)
 - Putty (windows)
 - Via remote desktop
 - X2GO (Windows/Mac/Linux)
 - Microsoft Remote Desktop (Windows/Mac/Linux)
- List of available iLab Machines
 - <https://report.cs.rutgers.edu/nagiosnotes/iLab-machines.html>

Connecting via SSH: Linux/Mac

- Open a new terminal
- Run “ssh <netid>@<ilabmachine>”
- Type in password to login



```
Davids-MacBook-Air:~ dj$ ssh djd240@ilab1.cs.rutgers.edu

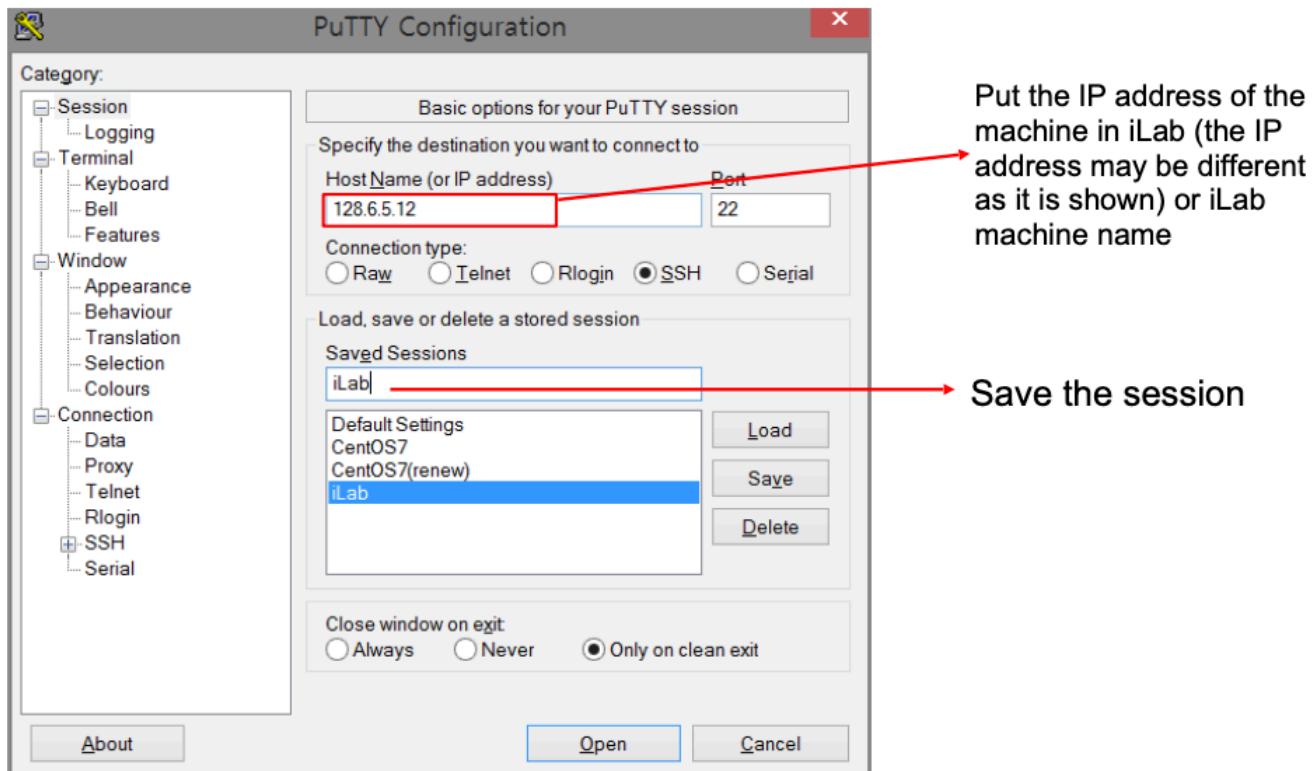
---IMPORTANT---
This ilab1.cs machine is using CS password server. If you have not done so,
"Activate" and "Set your CS password" at services.cs.rutgers.edu/accounts.

Password:
```

Guide: <https://www.howtogeek.com/311287/how-to-connect-to-an-ssh-server-from-windows-macos-or-linux/>

Connecting via SSH (Putty):Windows

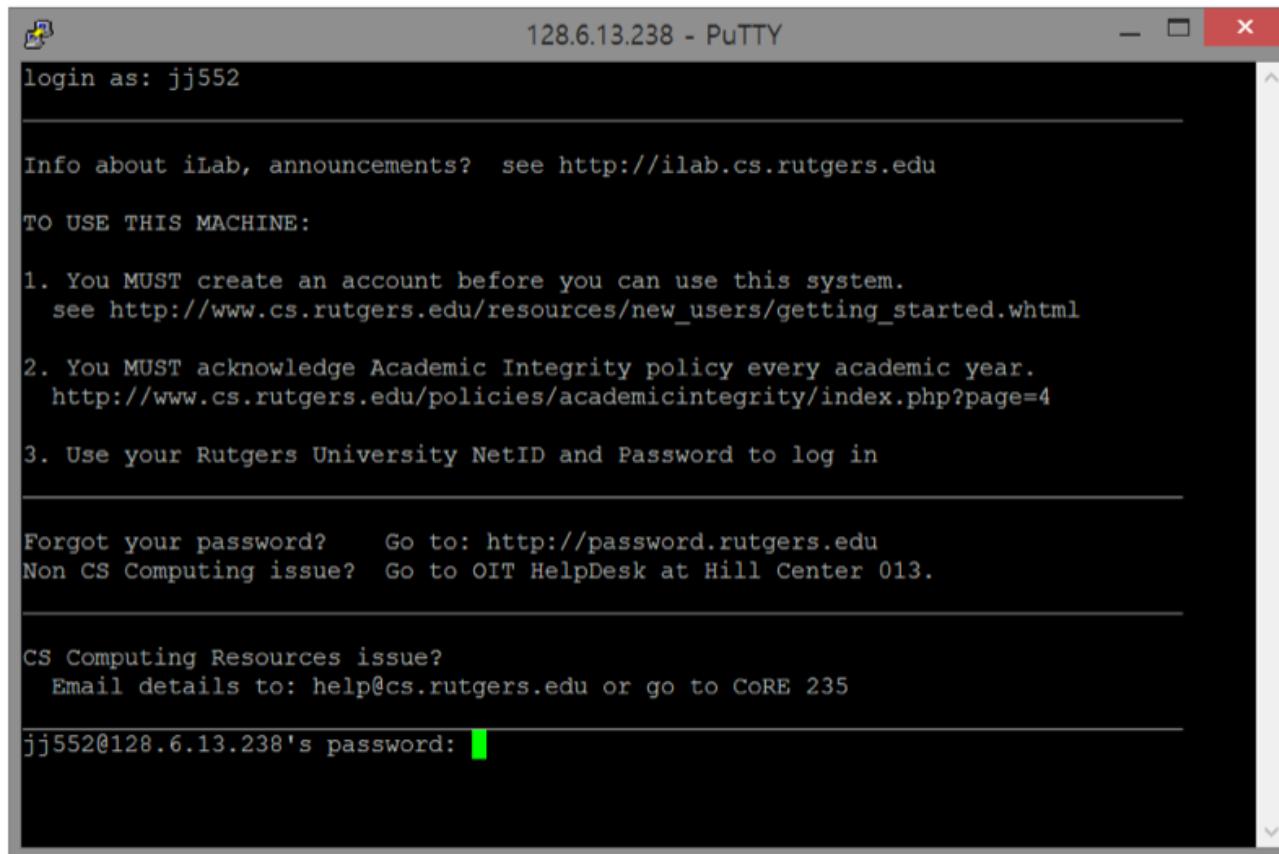
- Install Putty and configure Putty



Guide: <https://www.howtogeek.com/311287/how-to-connect-to-an-ssh-server-from-windows-macos-or-linux/>

Connecting via SSH (Putty):Windows

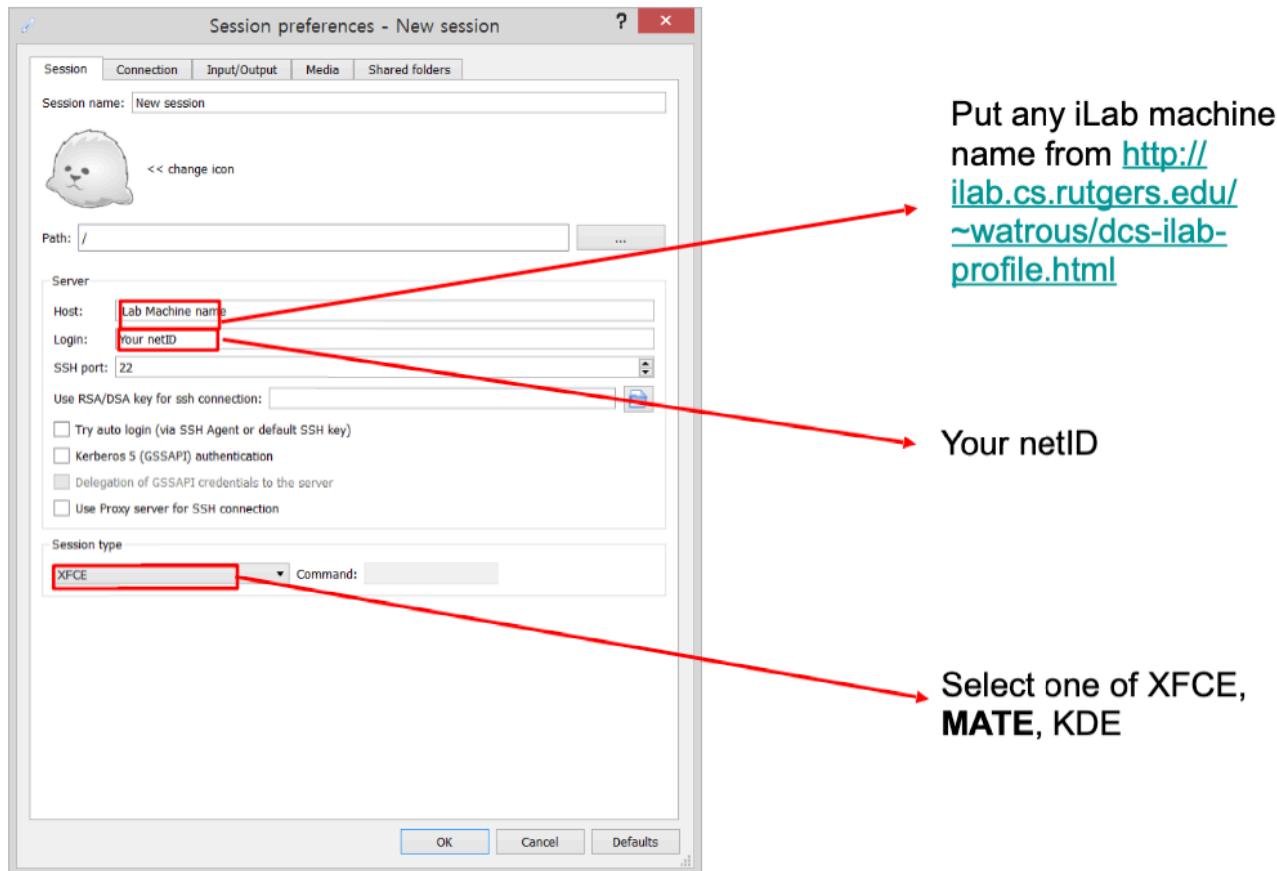
- Save session and open connection



Guide: <https://www.howtogeek.com/311287/how-to-connect-to-an-ssh-server-from-windows-macos-or-linux/>

Connecting to iLab Machines: X2GO Remote Desktop (Windows/Mac/Linux)

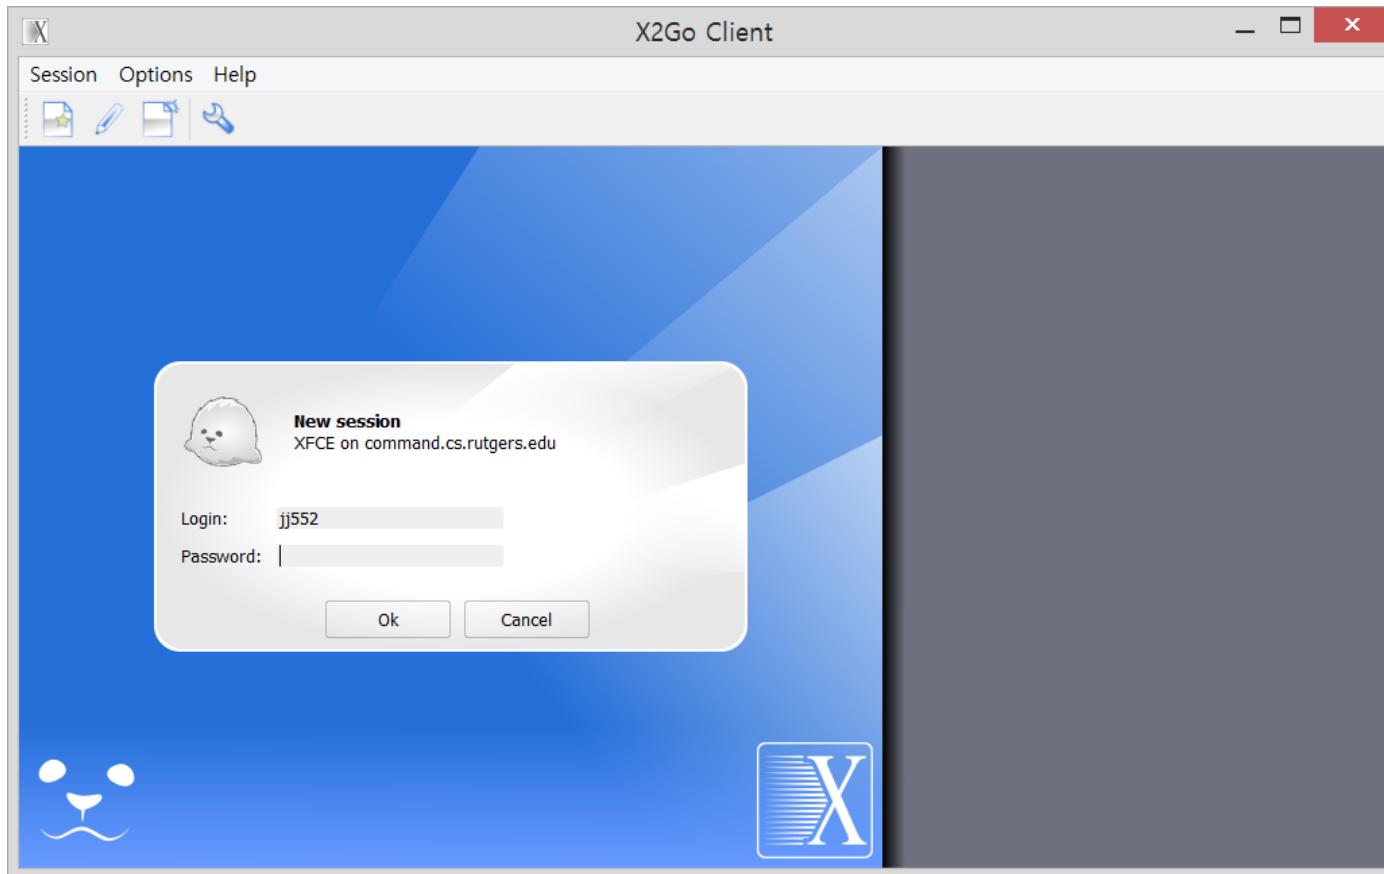
- Create new session



Guide: <https://resources.cs.rutgers.edu/docs/other/x2go/>

Connecting to iLab Machines: X2GO Remote Desktop (Windows/Mac/Linux)

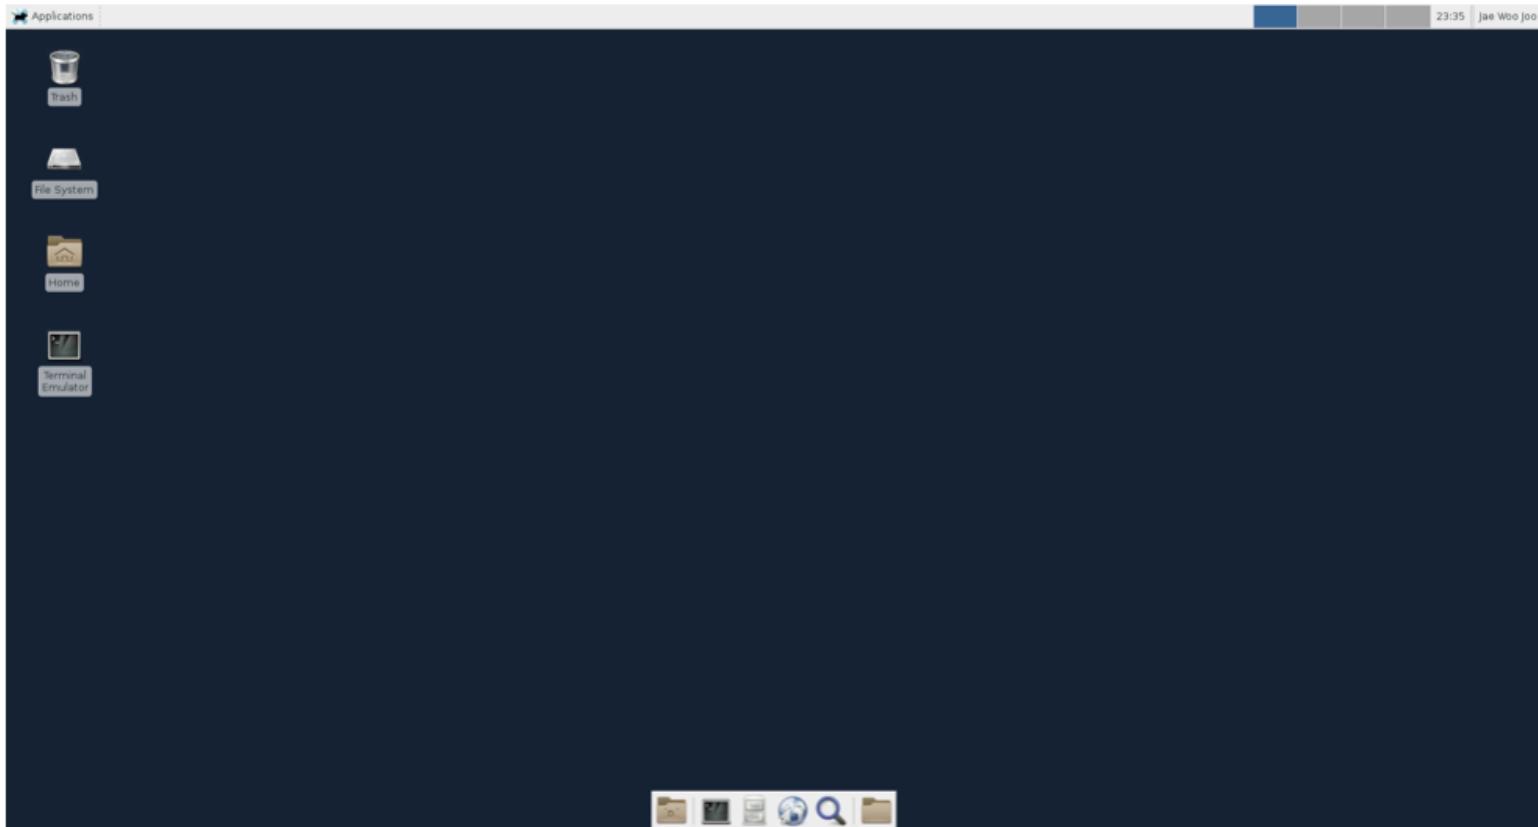
- Start session and login



Guide: <https://resources.cs.rutgers.edu/docs/other/x2go/>

Connecting to iLab Machines: X2GO Remote Desktop

- Start using Desktop Environment



Guide: <https://resources.cs.rutgers.edu/docs/other/x2go/>

Connecting to iLab Machines: Microsoft Remote Desktop (Windows/Mac/Linux)

- Necessary programs differ between systems
- Refer to guide linked below to see how to do it on your particular OS

Guide: <https://resources.cs.rutgers.edu/docs/computer-systems/windows-remote-desktop/>

Linux Shell Basics

- `pwd` : show name of current working directory
- `cd` : change directory
- `cd ..` : change to parent directory
- `ls` : list files
- `mkdir` : make a directory
- `rmdir` : remove a directory
- `touch` : create a file
- `cat` : print file contents
- `mv` : rename/move a file
- `rm` : remove a file
- `grep` : search texts in files

Special Directory Entries

- Every directory have two hidden directories:
 - “.” and “..”
- “.” refers to the current directory
- “..” refers to the parent directory
- Examples
 - cd .. (navigate to parent directory)
- <https://superuser.com/questions/37449/what-are-and-in-a-directory>

Linux Text Editors

- Vi / Vim
- Nano
- Gedit
- Emacs
-
- Pick anyone you would like

Vim Tutorial: <https://linuxconfig.org/vim-tutorial>

Nano Beginners Guide: <https://www.howtogeek.com/howto/42980/the-beginners-guide-to-nano-the-linux-command-line-text-editor/>

Transferring Files Between iLab and Local Machine

- Naïve way: Email files to yourself
- via scp (secure copy) on Mac/Linux
 - SCP tutorial: <https://kb.iu.edu/d/agye>
 - Examples:
 - (To iLab machine)
scp localfile.txt djd240@ilab.cs.rutgers.edu:folder/file.txt
 - (From iLab machine)
scp djd240@ilab.cs.rutgers.edu:folder/remotefile.txt file.txt
- via pscp (PuTTY secure copy) on Windows
 - pscp tutorial: <https://www.ssh.com/ssh/putty/putty-manuals/0.68/Chapter5.html>
 - Examples:
 - (To iLab machine)
pscp localfile.txt djd240@ilab.cs.rutgers.edu:folder/file.txt
 - (From iLab machine)
pscp djd240@ilab.cs.rutgers.edu:folder/remotefile.txt file.txt
- via SFTP client
 - drop and drag files between machines via GUI
 - WINSCP (Windows) <https://winscp.net/eng/index.php>
 - FileZilla Client (Windows/Mac/Linux) <https://filezilla-project.org/>